



Ontology (Knowledge Graph) Embeddings

Ernesto Jiménez-Ruiz

Lecturer in Artificial Intelligence

Before we start...

Students' module evaluation

- Your feedback is very important.
- Evaluations are anonymous.
- Access via e-mail from `evaluations@city.ac.uk` or from MyMoodle
- More information: `https://studenthub.city.ac.uk/
student-administration/online-module-evaluation-at-city`

Additional (potential) MSc Projects

DYAD: <https://www.dyad.net/>

- Healthcare domain.
- KG and graph machine learning.

Invited talks next week (15+5 min.)

- **Valentina Carapella**

- Data Scientist at Perspectum (<https://perspectum.com/>)
- *“KG Use Cases from Medical Imaging Science”*

Invited talks next week (15+5 min.)

- **Valentina Carapella**

- Data Scientist at Perspectum (<https://perspectum.com/>)
- *“KG Use Cases from Medical Imaging Science”*

- **Alexandra Shatova**

- MSc Data Science at City. Internship at DYAD (<https://www.dyad.net/>).
- *“Automated KG construction in the medical domain”*

Invited talks next week (15+5 min.)

- **Valentina Carapella**

- Data Scientist at Perspectum (<https://perspectum.com/>)
- *“KG Use Cases from Medical Imaging Science”*

- **Alexandra Shatova**

- MSc Data Science at City. Internship at DYAD (<https://www.dyad.net/>).
- *“Automated KG construction in the medical domain”*

- **Vicenzo Cutrona**

- PhD Student Università degli Studi di Milano - Bicocca.
- R&D OpenLab @ Corvallis SRL (<https://corvallis.it/>)
- *“Why semantic table understanding matters! Practical solutions to real-life problems”*

Where are we?

- ✓ Introduction.
- ✓ RDF-based knowledge graphs.
- ✓ SPARQL 1.0
- ✓ RDFS Semantics and RDF(S)-based knowledge graphs.
- ✓ OWL (2) ontology language. Focus on modelling.
- ✓ Application to Data Science.
- ✓ OWL 2 Profiles, SPARQL 1.1 and Entailment Regimes
- ✓ Ontology (Knowledge Graph) Alignment

Where are we?

- ✓ Introduction.
- ✓ RDF-based knowledge graphs.
- ✓ SPARQL 1.0
- ✓ RDFS Semantics and RDF(S)-based knowledge graphs.
- ✓ OWL (2) ontology language. Focus on modelling.
- ✓ Application to Data Science.
- ✓ OWL 2 Profiles, SPARQL 1.1 and Entailment Regimes
- ✓ Ontology (Knowledge Graph) Alignment

9. **Ontology (Knowledge Graph) Embeddings (today).**

Where are we?

- ✓ Introduction.
- ✓ RDF-based knowledge graphs.
- ✓ SPARQL 1.0
- ✓ RDFS Semantics and RDF(S)-based knowledge graphs.
- ✓ OWL (2) ontology language. Focus on modelling.
- ✓ Application to Data Science.
- ✓ OWL 2 Profiles, SPARQL 1.1 and Entailment Regimes
- ✓ Ontology (Knowledge Graph) Alignment

9. **Ontology (Knowledge Graph) Embeddings (today).**

10. Graph Database Solutions and [Invited Talks](#) (March 31).

Hybrid Learning and Reasoning Systems

Motivation:

- Need of **richer AI** systems, *i.e.*, **semantically sound, explainable, and reliable**.

Gary Marcus. The Next Decade in AI: Four Steps Towards Robust Artificial Intelligence. CoRR abs/2002.06177 (2020)

Motivation:

- Need of **richer AI** systems, *i.e.*, **semantically sound, explainable, and reliable**.
- Impressive results in Deep Learning but require **large datasets** and **lack explanation**.

Gary Marcus. The Next Decade in AI: Four Steps Towards Robust Artificial Intelligence. CoRR abs/2002.06177 (2020)

Motivation:

- Need of **richer AI** systems, *i.e.*, **semantically sound, explainable, and reliable**.
- Impressive results in Deep Learning but require **large datasets** and **lack explanation**.
- Limitations of KR systems: **maintenance** and **flexibility** in the inference. *e.g.*, Does $C(a)$ hold if?
 - A and $(R \text{ some } B)$ subClassOf C . $A(a)$, $B'(b)$, and $R(a, b)$.

Gary Marcus. The Next Decade in AI: Four Steps Towards Robust Artificial Intelligence. CoRR abs/2002.06177 (2020)

Motivation:

- Need of **richer AI** systems, *i.e.*, **semantically sound, explainable, and reliable**.
- Impressive results in Deep Learning but require **large datasets** and **lack explanation**.
- Limitations of KR systems: **maintenance** and **flexibility** in the inference. *e.g.*, Does $C(a)$ hold if?
 - A and $(R \text{ some } B)$ subClassOf C . $A(a)$, $B'(b)$, and $R(a, b)$.
- **Solution?** Hybrid Learning and Reasoning Systems.

Gary Marcus. The Next Decade in AI: Four Steps Towards Robust Artificial Intelligence. CoRR abs/2002.06177 (2020)

Hybrid Learning and Reasoning Systems

- Unification of:
 - **statistical** (data-driven) and
 - **symbolic** (knowledge-driven) methods

† Michael van Bekkum et al. Modular Design Patterns for Hybrid Learning and Reasoning Systems: a taxonomy, patterns and use cases. CoRR abs/2102.11965. Under review (2021)

Hybrid Learning and Reasoning Systems

- Unification of:
 - **statistical** (data-driven) and
 - **symbolic** (knowledge-driven) methods
- Overview of **patterns** for hybrid systems. †

† Michael van Bekkum et al. Modular Design Patterns for Hybrid Learning and Reasoning Systems: a taxonomy, patterns and use cases. CoRR abs/2102.11965. Under review (2021)

Hybrid Learning and Reasoning Systems

- Unification of:
 - **statistical** (data-driven) and
 - **symbolic** (knowledge-driven) methods
- Overview of **patterns** for hybrid systems. †
 - Focus on **Ontology** (knowledge graph) **embeddings** as a component for an hybrid system (*e.g.*, OWL2Vec*).

† Michael van Bekkum et al. Modular Design Patterns for Hybrid Learning and Reasoning Systems: a taxonomy, patterns and use cases. CoRR abs/2102.11965. Under review (2021)

Generic Patterns for Hybrid Systems

Focus on Knowledge Graph Embeddings

Introduction: Models

- **Models** are descriptions of entities and their relationships.

Introduction: Models

- **Models** are descriptions of entities and their relationships.
- **Statistical models** represent dependencies between statistical variables.
 - *e.g.*, Examples are (Deep) Neural Networks, Bayesian Networks and Markov Models.

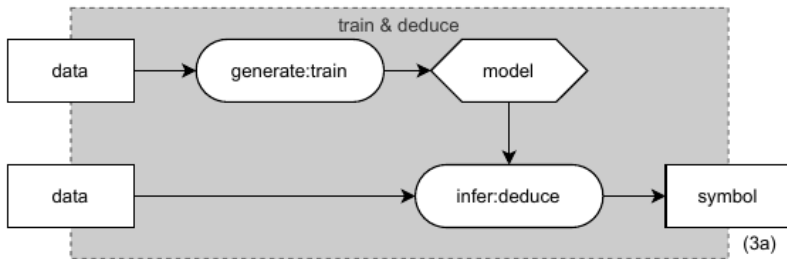
Introduction: Models

- **Models** are descriptions of entities and their relationships.
- **Statistical models** represent dependencies between statistical variables.
 - *e.g.*, Examples are (Deep) Neural Networks, Bayesian Networks and Markov Models.
- **Semantic models** represent the implicit meaning and relationships of symbols.
 - *e.g.*, Ontologies, Knowledge Graphs, Rule-based models.

Introduction: Models

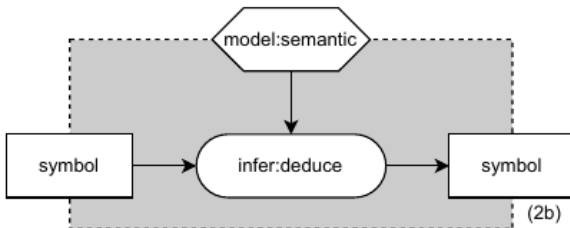
- **Models** are descriptions of entities and their relationships.
- **Statistical models** represent dependencies between statistical variables.
 - *e.g.*, Examples are (Deep) Neural Networks, Bayesian Networks and Markov Models.
- **Semantic models** represent the implicit meaning and relationships of symbols.
 - *e.g.*, Ontologies, Knowledge Graphs, Rule-based models.
- **Hybrid models** combine both.

Machine learning pattern (non hybrid)



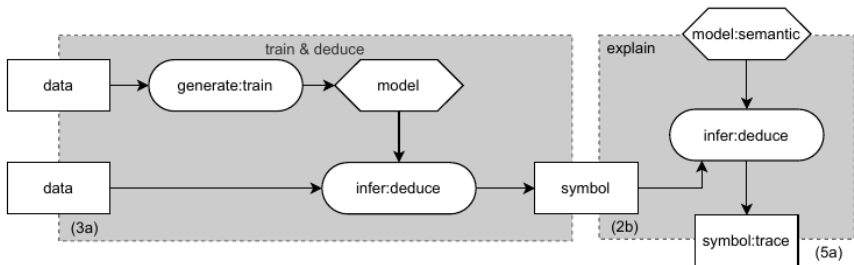
For example, image classification as in the <http://www.image-net.org/> challenge (symbol = label from WordNet).

Semantic model pattern (non hybrid)



- Standard reasoning (*e.g.*, classification, class membership).
- Rule-mining and ontology learning based on symbolic data (*i.e.*, ABox).

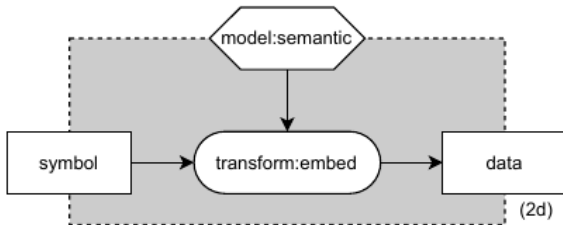
Explainability



The semantic model explains/interprets the prediction.

Explaining Trained Neural Networks with Semantic Web Technologies: First Steps. NeSy 2017.
Human-driven FOL explanations of deep learning. IJCAI 2020

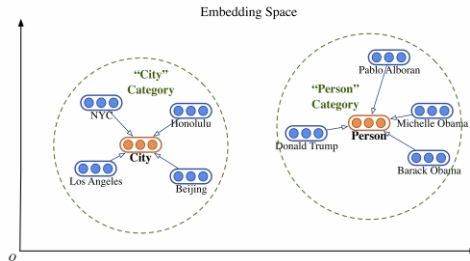
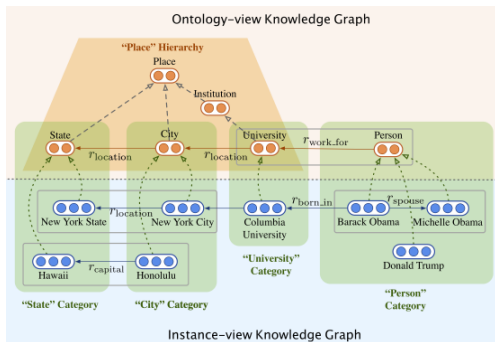
Knowledge graph embeddings



Symbols are transformed into vectors (*e.g.*, OWL2Vec)

Knowledge Graph Embedding: A Survey of Approaches and Applications. TKDE 2017
OWL2Vec*: Embedding of OWL Ontologies. CoRR abs/2009.14654 (2020)

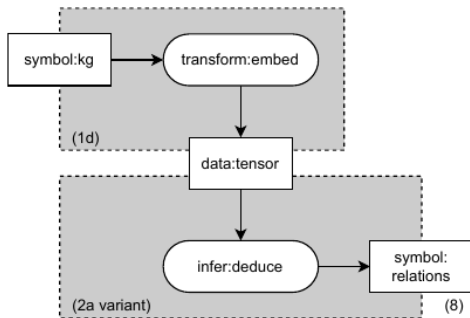
Knowledge graph embeddings (example)



KG Embedding Systems exploit the neighbourhood of an entity to calculate its vector.

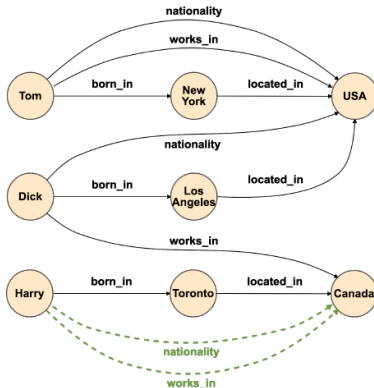
Example from: Universal Representation Learning of Knowledge Bases by Jointly Embedding Instances and Ontological Concepts. KDD 2019.

Knowledge Graph Embeddings: Link prediction



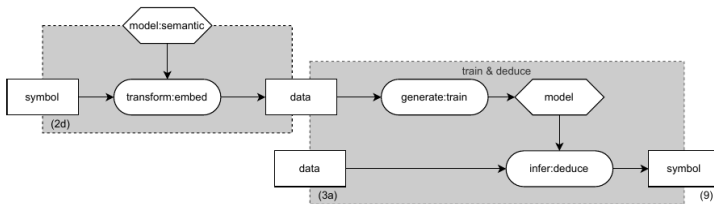
- Plausability of a triple <subject predicate object> given a scoring function.

Knowledge Graph Embeddings: Link prediction (example)



Example from: Knowledge Graph Embedding for Link Prediction: A Comparative Analysis. TKDD 2021

Learning with (knowledge) embeddings

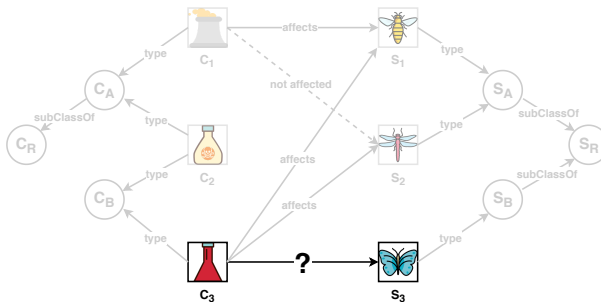


- Applying knowledge graph embeddings in a subsequent classification step.
- Graph Neural Networks over the KG structure.
- Key for zero-shot learning approaches

Prediction of Adverse Biological Effects of Chemicals Using Knowledge Graph Embeddings. Under review. 2021.
A Comprehensive Survey on Graph Neural Network. IEEE Transactions on Neural Networks and Learning Systems 2019.
Knowledge-aware Zero-Shot Learning: Survey and Perspective. arXiv:2103.00070. 2021

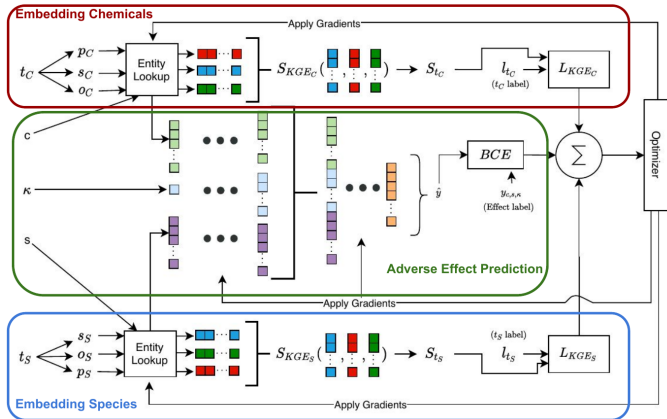
Learning with (knowledge) embeddings (example)

- **Prediction of adverse biological effects** of chemicals via KG embeddings.



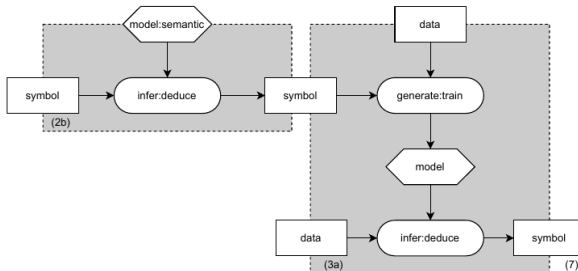
Prediction of Adverse Biological Effects of Chemicals Using Knowledge Graph Embeddings. Under review. 2021.

Learning with (knowledge) embeddings (example)



Prediction of Adverse Biological Effects of Chemicals Using Knowledge Graph Embeddings. Under review. 2021.

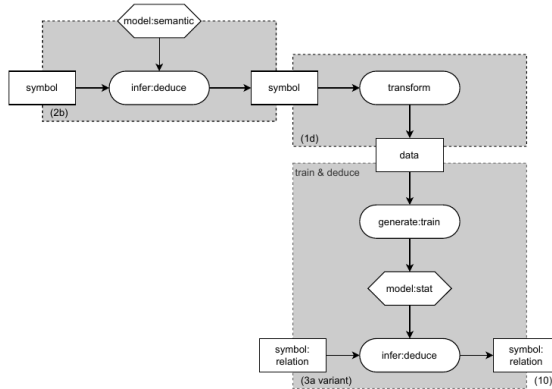
Learning with prior knowledge



- Domain knowledge (e.g., a KG) used to constraint search space during training.
- **Semantic loss function:** impact of the violation of the symbolic knowledge.

A semantic loss function for deep learning with symbolic knowledge. ICML 2018
Logic Tensor Networks. <https://github.com/logictensornetworks/logictensornetworks>

Learning to reason

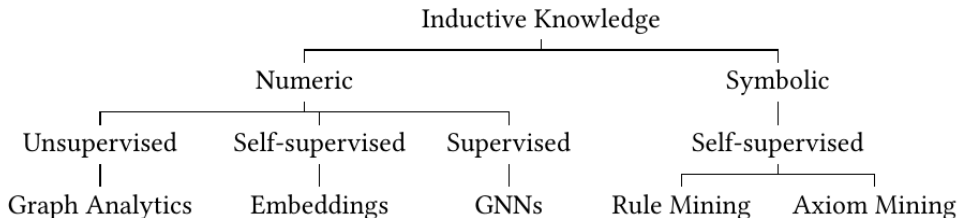


Ontology Reasoning with Deep Neural Networks. JAIR 2021

Inductive Techniques for Knowledge Graphs

Focus on Knowledge Graph Embeddings

Inductive techniques for knowledge graphs



- Input: A Knowledge Graph (symbolic)
- Output: Numeric or Symbolic

Knowledge Graphs. arXiv:2003.02320. 2021

Graph analytics (unsupervised)

Exploit techniques from **graph theory** and **network analysis**. *e.g.*,:

- **Centrality**: the most important nodes (*i.e.*, concepts, instances) or edges (*i.e.*, properties) of a graph.
- **Community detection**: subgraphs that are densely connected.
- **Connectivity**: how well-connected are the nodes of a graph to identify isolated nodes or subgraphs.

Graph analytics (unsupervised)

Exploit techniques from **graph theory** and **network analysis**. *e.g.*,:

- **Centrality**: the most important nodes (*i.e.*, concepts, instances) or edges (*i.e.*, properties) of a graph.
- **Community detection**: subgraphs that are densely connected.
- **Connectivity**: how well-connected are the nodes of a graph to identify isolated nodes or subgraphs.
- **Path finding**: all possible paths between two nodes.
- **Node similarity**: based on their connection to other nodes (*e.g.*, random-walks techniques).

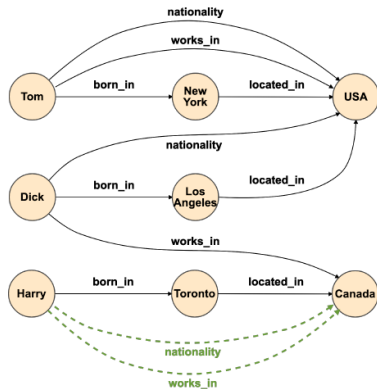
Graph Neural Networks (supervised)

- Machine learning models for **graph-structured data**.
- The **neural network** is **based on the shape and connections** of the (knowledge) graph
- **End-to-end supervised learning** (*e.g.*, classification).
- Can be used to classify nodes or the graph itself.

A Comprehensive Survey on Graph Neural Network. IEEE Transactions on Neural Networks and Learning Systems 2019.

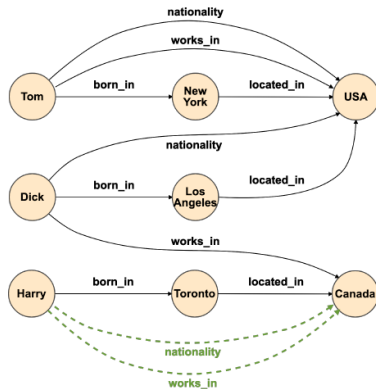
Symbolic learning (self-supervised)

- Identifies patterns in the data



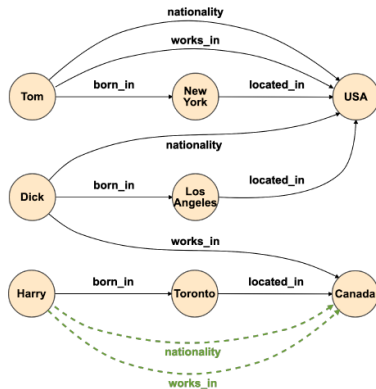
Symbolic learning (self-supervised)

- Identifies patterns in the data
- Learn hypotheses in a symbolic (logical) language. For example:
 - As a rule: $\text{nationality}(x,z) :- \text{born_in}(x,y) \wedge \text{located_in}(y,z)$
 - As an OWL 2 axioms:
 $\text{born_in} \circ \text{located_in}$
 $\text{SubPropertyOf: nationality}$

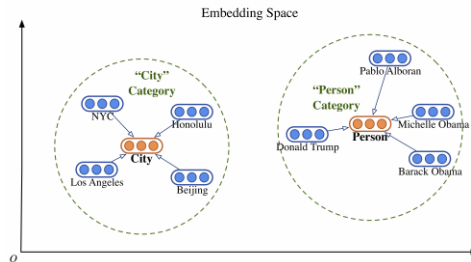
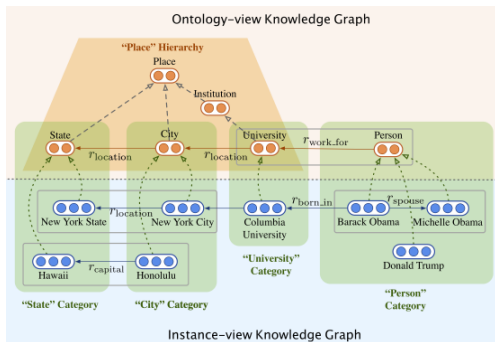


Symbolic learning (self-supervised)

- Identifies patterns in the data
- Learn hypotheses in a symbolic (logical) language. For example:
 - As a rule: $\text{nationality}(x,z) :- \text{born_in}(x,y) \wedge \text{located_in}(y,z)$
 - As an OWL 2 axioms:
 $\text{born_in} \circ \text{located_in}$
 $\text{SubPropertyOf: nationality}$
- Can help explaining/interpret (link) predictions: *e.g.*, why $\text{nationality}(\text{Hernry}, \text{Canada})?$



Knowledge graph embeddings (self-supervised)



Example from: Universal Representation Learning of Knowledge Bases by Jointly Embedding Instances and Ontological Concepts. KDD 2019.

Knowledge graph embeddings (self-supervised)

KGE approaches (excluding those based on language models) typically:

- Receive as input a set of **positive** (the ones in the KG) and **negative triples**.
- Include a **scoring function** that accepts as input the embedding of the elements of a triple (there is an initialization step).

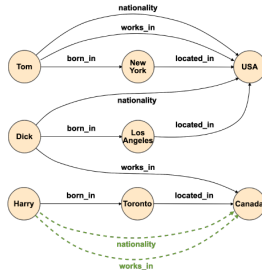
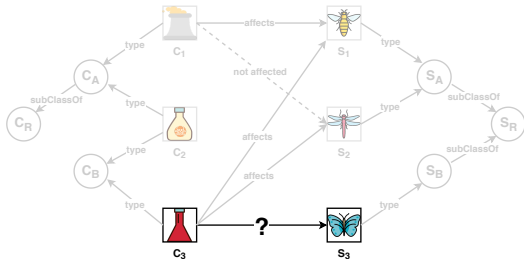
Knowledge graph embeddings (self-supervised)

KGE approaches (excluding those based on language models) typically:

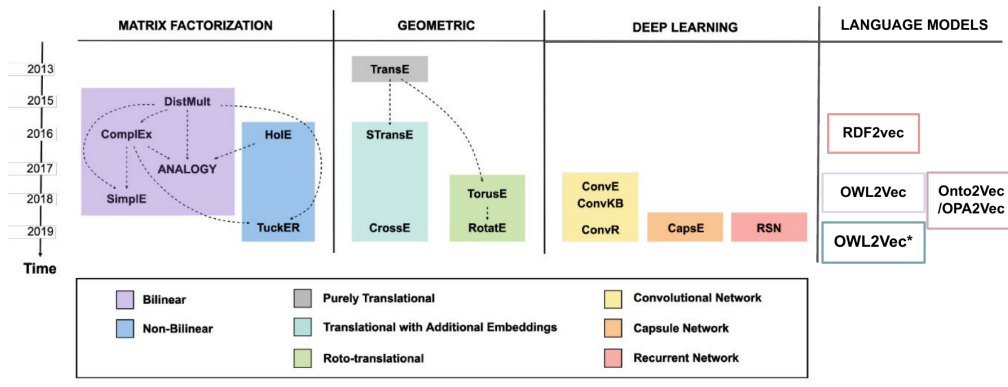
- Receive as input a set of **positive** (the ones in the KG) and **negative triples**.
- Include a **scoring function** that accepts as input the embedding of the elements of a triple (there is an initialization step).
- Learn embedding so that the score for positive triples is maximized while the score for negative triples is minimized (*i.e.*, **loss function**).
- Compute **similar vectors** for similar nodes (*i.e.*, concepts/instances) and edges (*i.e.*, properties).

Knowledge graph embeddings (applications)

- The computed embedding can be used in a **downstream machine learning task** (e.g., prediction of adverse effect chemical-species).
- The scoring function can be used to evaluate the plausibility of a triple for **link prediction** or **KG completion**.



Knowledge graph embeddings (overview of approaches)



Incomplete list of approaches, adapted from: Knowledge Graph Embedding for Link Prediction: A Comparative Analysis. TKDD 2021

Knowledge graph embeddings (overview of approaches)

- **Translational models:** translate subject entities to object entities via the predicate/relation in the low-dimensional space.

Knowledge graph embeddings (overview of approaches)

- **Translational models:** translate subject entities to object entities via the predicate/relation in the low-dimensional space.
- **Tensor decomposition models:** represent the KG into a one-hot 3-order tensor and apply matrix factorisation/decomposition models to learn entity vectors.

Knowledge graph embeddings (overview of approaches)

- **Translational models:** translate subject entities to object entities via the predicate/relation in the low-dimensional space.
- **Tensor decomposition models:** represent the KG into a one-hot 3-order tensor and apply matrix factorisation/decomposition models to learn entity vectors.
- **Neural models:** unlike previous models they learn embeddings with non-linear scoring functions via a neural model.

Knowledge graph embeddings (overview of approaches)

- **Translational models**: translate subject entities to object entities via the predicate/relation in the low-dimensional space.
- **Tensor decomposition models**: represent the KG into a one-hot 3-order tensor and apply matrix factorisation/decomposition models to learn entity vectors.
- **Neural models**: unlike previous models they learn embeddings with non-linear scoring functions via a neural model.
- **Language models**: perform random walks over the KG to create a document of sentences and leverage existing language models (*e.g.*, word embedding) to learn vectors for each KG entity.

Knowledge graph embeddings (implementations)

- PyKEEN (Python KnowlEdge EmbeddiNgs) with PyTorch:
<https://pykeen.github.io/>
- Open Knowledge Embedding implemented with PyTorch:
<https://github.com/thunlp/OpenKE>
- Knowledge Embedding implemented with Keras:
<https://github.com/NIVA-Knowledge-Graph/KGE-Keras>
- jRDF2Vec: <https://github.com/dwslab/jRDF2Vec>
- pyRDF2Vec: <https://github.com/IBCNServices/pyRDF2Vec>
- OWL2Vec* (python): <https://github.com/KRR-Oxford/OWL2Vec-Star>

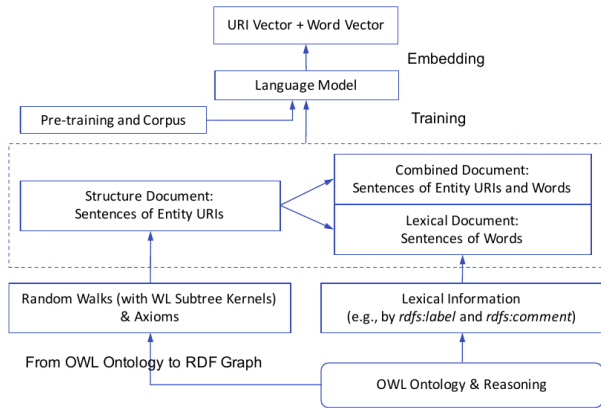
Knowledge graph embeddings (pre-trained)

- OpenKE:
<http://139.129.163.161/index/toolkits#pretrained-embeddings>
- KGvec2go: <http://www.kgvec2go.org/>
- Drug-drug interaction:
<https://github.com/rcelebi/GraphEmbedding4DDI/>

Embedding ontologies with OWL2Vec*

OWL2Vec* Overview

- **projects** the ontology into a graph,
- **walks** the graph,
- creates a **corpus of sentences** according to the walking strategies, and
- generates **embeddings** from that corpus.



OWL2Vec*: Embedding of OWL Ontologies. CoRR abs/2009.14654 (2020)

OWL2Vec*: ontology projection

Approximation of an OWL 2 ontology into an RDF graph.

Axiom of Condition 1	Axiom or Triple(s) of Condition 2	Projected Triple(s)
$A \sqsubseteq \Box r.D$ or $\Box r.D \sqsubseteq A$	$D \equiv B \mid B_1 \sqcup \dots \sqcup B_n \mid B_1 \sqcap \dots \sqcap B_n$	$\langle A, r, B \rangle$ or $\langle A, r, B_i \rangle$ for $i \in 1, \dots, n$
$\exists r. \top \sqsubseteq A$ (domain)	$\top \sqsubseteq \forall r. B$ (range)	
$A \sqsubseteq \exists r. \{b\}$	$B(b)$	
$r \sqsubseteq r'$	$\langle A, r', B \rangle$ has been projected	
$r' \equiv r^-$	$\langle B, r', A \rangle$ has been projected	
$s_1 \circ \dots \circ s_n \sqsubseteq r$	$\langle A, s_1, C_1 \rangle \dots \langle C_n, s_n, B \rangle$ have been projected	
$B \sqsubseteq A$	–	$\langle B, rdfs:subClassOf, A \rangle$ $\langle A, rdfs:subClassOf^-, B \rangle$
$A(a)$	–	$\langle a, rdf:type, A \rangle$ $\langle A, rdf:type^-, a \rangle$
$r(a, b)$	–	$\langle a, r, b \rangle$

\Box is one of: $\geq, \leq, =, \exists, \forall$. A, B, B_i and C_i are atomic concepts (classes), s_i, r and r' are roles (object properties), r^- is the inverse of a relation r , a and b are individuals, \top is the top concept.

OWL2Vec*: sentence generation via random walks

Strategies:

- Random walks
- Weisfeiler Lehman (WL) kernel, which assign identifiers to subgraphs and includes them into the walk.

Structure Document Sentences

(vc:Beer, rdf:type, vc:FOOD-4001, vc:hasNutrient, vc:VitaminC_1000)

Lexical Document Sentences

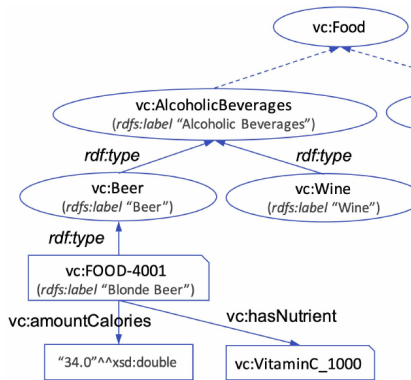
(“beer”, “type”, “blonde”, “beer”, “has”, “nutrient”, “vitamin”, “c”)

Combined Document Sentences

(vc:FOOD-4001, “has”, “nutrient”, “vitamin”, “c”)

OR

(“blonde”, “beer”, “has”, “nutrient”, vc:VitaminC_1000)



OWL2Vec*: language model and embeddings

- OWL2Vec relies on the **Word2vec** as neural **language model**.
- Word2vec learns **embeddings** for all the elements in the documents (*i.e.*, both **words** and **URIs**)

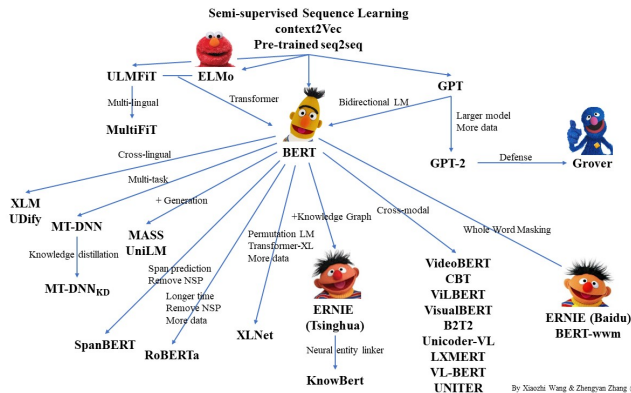
OWL2Vec*: language model and embeddings

- OWL2Vec relies on the **Word2vec** as neural **language model**.
- Word2vec learns **embeddings** for all the elements in the documents (*i.e.*, both **words** and **URIs**)
- The embeddings of the ontology entities can be calculated via their **URI embedding** or via the **word embeddings** of their labels.
 - The URI `vc:FOOD-4001` (Blonde Beer) has a vector.
 - As well as the words ‘blonde’ and ‘beer’.

OWL2Vec*: language model (future)

Other language models could be used in OWL2Vec*:

<https://github.com/thunlp/PLMpapers>



By Xiaochi Wang & Zhengyan Zhang @THUNLP

OWL2Vec*: applications

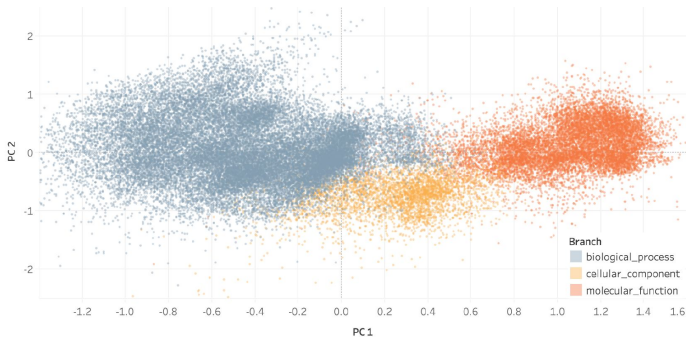
- **Class subsumption** and **class membership predictions** (as in OWL2Vec* paper)
- Embedding of chemicals and species to **predict adverse effects**.
- **Ontology alignment** (Samsung UK project with food ontologies) †
- **Ontology clustering** in life sciences ontologies to be applied in an Information Retrieval task. ‡

† J. Chen et al. Augmenting Ontology Alignment by Semantic Embedding and Distant Supervision. ESWC 2021

‡ A. Ritchie. Ontology Clustering with OWL2Vec*. Submitted 2021.

OWL2Vec*: clustering

Embedding of the Gene Ontology and its 3 branches: biological process, cellular component, and molecular function



A. Ritchie. Ontology Clustering with OWL2Vec*. Submitted 2021.

Acknowledgements

- OWL2Vec* developers and collaborators.
- Specially **Jiaoyan Chen**, University of Oxford

Laboratory Session

OWL2Vec* in practice

- Execute OWL2Vec* over the `Pizza` and `FoodOn` ontologies.
- Compute similarity among words and entities.
- Perform clustering and visualize results.